

Your Security Future Demands a Fundamental Rethink

Your Team Already Had the Data

Every time there is a breach, your security team sets out to determine exactly what happened. They reconstruct how attackers compromised your systems, map which accounts were involved, and trace what data was stolen. Then, in every postmortem, the uncomfortable truth reveals itself: You had the information all along.

The question that should keep you awake is: If you have the information to find out what happened after the fact, why can't you do something to prevent or stop the incident before it happens?

This is a design problem, not a technology problem.

The Query Hamster Wheel

Your analysts start each day by opening seven different security consoles. They check the security information and event management (SIEM) platform for overnight alerts and then switch to the endpoint protection platform. Next, they open the cloud security tool and log into the identity management system. Each switch costs mental energy. Each query requires a different syntax. Each investigation demands manual correlation across disconnected systems.

When an authentication failure alert fires, your analyst must become a detective across fragmented evidence. Three hours later, your analyst determines this was a misconfigured application, not an attack. Meanwhile, the real threat moved laterally across your network undetected because it generated alerts in three different systems that nobody connected.

This is not how defense should work. They act more as tool administrators, not defenders.

Why Traditional Approaches Fail

The cybersecurity industry responds to this challenge with a simple answer: Make search faster. But faster search misses the fundamental question. Why are your analysts querying all day to begin with?

If your security platform understood your environment, it would surface threats automatically. Your team would investigate incidents, not chase alerts. They would make strategic decisions, not run endless queries.

Instead, you are stuck in a model where more tools supposedly equal better coverage, more data supposedly equals more visibility, and more automation supposedly equals transformation.

None of this is true.

More tools create more complexity. Your team now maintains dozens of security products, each requiring dedicated expertise. Each integration point creates brittleness. Each vendor relationship demands management overhead. You spend millions of dollars on technology while your analysts might spend most of their time being tool administrators.

More data creates a digestion problem, not a detection advantage. You ingest terabytes of security telemetry daily but cannot make sense of it fast enough. The data exists. The visibility is there. The problem is connecting disconnected dots across siloed systems before attackers complete their objectives.

Automation does not increase visibility, but rather only introduces discipline into manual processes, because you can't automate a process that is broken. For example, you might automate alert enrichment, but the alerts remain disconnected. While you might orchestrate response workflows, the correlation stays manual. And, although you build playbooks for known scenarios, novel attacks slip through undetected.

Traditional security operations optimize for yesterday's threats while attackers exploit tomorrow's techniques.

What Could Be: The Autonomous SOC Vision

Imagine your security operations center tomorrow morning. Your analysts arrive to find the overnight incidents already correlated, enriched, and prioritized. Low-confidence alerts from various systems have been stitched into high-confidence incidents that show complete attack timelines. Routine threats have been contained automatically. Your team focuses immediately on the three incidents that require human judgment.

An authentication anomaly appears not as a single failed login, but as a fully contextualized incident. The platform shows the user's normal behavior patterns: recent changes to access rights, other activities from the source location, similar patterns across your environment, relevant threat intelligence, business impact assessment, and recommended response actions.

Your analyst makes a decision in minutes instead of hours because the investigation work is already done. The platform performed the correlation, and the enrichment happened automatically. The business context is clear.

This is the autonomous SOC. Technology handles what technology does well while humans focus on what requires human judgment.

AI Agents: Security That Reasons and Acts

The autonomous SOC requires more than faster queries or prebuilt playbooks. It needs technology that can reason through novel scenarios, adapt to changing contexts, and take intelligent actions without waiting for human-defined logic.

This is where agentic AI transforms security operations.

Traditional automation follows predetermined paths. If X happens, then do Y. These playbooks work well for known scenarios, but attackers don't follow scripts. They adapt. They pivot.

Agentic AI operates differently. Rather than executing fixed sequences, AI agents analyze the complete context of a security incident and determine the appropriate response dynamically. They understand the business impact of their actions, recognize when situations fall outside normal patterns, and escalate intelligently only when they need human judgment.

When your endpoint agent detects suspicious process execution, an AI agent doesn't just follow a playbook. It examines the complete picture: the user's typical behavior patterns, the application's normal operation, similar activities across your environment, current threat intelligence, and your organization's risk tolerance. Based on this analysis, it determines whether to contain immediately, monitor closely, or escalate for review.

The agent's decision considers factors no static playbook could anticipate. Is this a developer testing new code or an attacker establishing persistence? Is the affected system critical to operations right now? Are there similar patterns emerging elsewhere that suggest a coordinated campaign? The agent synthesizes these inputs and acts accordingly.

Your security operations gain something unprecedented—technology that handles both the routine scenarios and the novel ones that sit between your playbooks. The system fills the gaps where attackers traditionally hide.

The Palo Alto Networks Vision: Security First, Data Second

Most security platforms are built backwards. They start with data collection and then try to figure out what to do with it. The philosophy is simple, if flawed: Ingest as much as possible and then determine the value later.

We believe you must rebuild security operations from the opposite direction. Start with the security outcomes you need to achieve. Determine what data and analytics deliver those outcomes. Build the intelligence first, and then design the data architecture to support that intelligence.

Security first, data second.

This inversion changes everything about how security operations work.

Why We Think This: The Practitioner Perspective

We built our vision of autonomous security operations based on two decades of defending enterprise environments and responding to sophisticated attacks. Our perspective comes from actual security practice, not theoretical models.

When Unit 42® responds to the world's most consequential breaches, our team deploys our security operations platform during the investigation. They use it to reconstruct attack timelines and to identify all compromised systems. And, they rely on it to ensure containment. Every engagement teaches us what security operations must become.

We protect the world's largest cybersecurity company with our own transformed SOC. We have helped hundreds of the world's largest enterprises transform their SOC as a partner through Unit 42 SOC Assessments and by setting up their own transformed SOC with Cortex XSIAM®.

We learned that the traditional SOC model creates three fundamental failures.

First, it treats all security data equally when some data matters far more than others for specific threats. Detecting credential theft requires different telemetry than identifying data exfiltration. Generic log collection misses this reality.

Second, it assumes analysts can manually correlate events when modern attacks span too many systems for human correlation to scale. Multistage campaigns that touch endpoints, networks, clouds, and identities generate dozens of disconnected alerts. Manual correlation cannot keep pace.

Third, it measures activity instead of outcomes. Mean time to detect (MTTD) tells you nothing about prevention effectiveness. Alert closure rates provide no indication of whether you eliminated root causes. Speed metrics optimize for reaction when you need to focus on resilience.

These failures are not vendor-specific. They are model-specific. No amount of query optimization fixes a broken operational approach.

Three Foundational Shifts

The SOC transformation requires fundamentally rethinking how security operations work.

Shift One: From Collection to Intelligence

Traditional security operations collect data and then expect you to build the intelligence layer yourself. You write the detection rules, create the correlation logic, and develop the analytics. The platform provides storage and search, but the security outcomes are your responsibility.

The autonomous approach inverts this model. First, you build the intelligence layer, with opinionated security knowledge embedded in the platform from day one.

You start by understanding the security outcomes that matter. Rather than building detection logic for each outcome from scratch, you leverage prebuilt analytics designed by security practitioners who have defended thousands of enterprises.

The problem was never data ingestion. It was data digestion. The real challenge begins after collection where you transform raw telemetry into intelligence that can stop threats. An intelligence-first data architecture processes, enriches, and structures security data as it arrives, turning raw events into contextualized intelligence before storage, not after.

Collect once, analyze infinitely. As you develop new security modules and capabilities, they leverage the existing data lake without requiring re-collection. Your investment in data collection compounds over time as new analytics continuously extract additional value from the same telemetry.

The result is that your platform delivers answers instead of waiting for questions. It neutralizes threats instantly through industry-leading automation that runs natively across the entire platform, optimizing your security outcomes and driving significant time savings and cost reductions.

Shift Two: From Alerts to Incidents

Your analysts today receive alerts—failed authentication, suspicious process execution, and unusual network connection. Each alert is a data point requiring investigation to determine its relevance.

You need incidents instead of alerts.

When suspicious authentication activity emerges, the platform automatically correlates it with endpoint behavior, network connections, data access patterns, and business context. Multiple low-confidence signals become a single high-confidence incident showing the complete picture.

You need actionable intelligence, not faster alert processing.

More importantly, you should aim to stop attacks in progress, not just after damage occurs. When the platform detects credential compromise patterns early, it can automatically enforce additional verification, restrict access to critical systems, or initiate containment before attackers achieve their objectives. Prevention happens at machine speed, while the investigation context builds in parallel so, when your analyst reviews the incident, they see both what was prevented and why.

Shift Three: From Reactive Response to Continuous Prevention

Traditional security operations react to threats. The autonomous approach prevents them.

This starts with learning, which extends beyond your organization. Threat intelligence from tens of thousands of enterprises continuously updates detection models. Attack techniques identified anywhere become protection everywhere. Your security posture improves daily without your team writing detection rules. Consider the following examples.

Example 1: Phishing Investigation Pattern

Your analyst investigates a phishing alert. They check whether other users received similar emails, query for clicked links, search for credential submissions, and verify whether MFA stopped the attack. After seeing this investigation sequence multiple times, the platform suggests automating it. Now, phishing alerts arrive with this analysis already complete.

Example 2: Suspicious Login Investigation

An analyst investigates failed authentication attempts. They always check the user's normal login locations, verify recent password changes, look for successful logins from the same source, and review access to sensitive resources. The platform recognizes this pattern and executes these queries automatically, presenting the complete context rather than just the initial alert.

Example 3: Automated Response Playbook

When analysts confirm a credential compromise, they consistently disable the account, revoke active sessions, reset passwords, and notify the user's manager. After observing this response pattern, the platform offers to automate these actions. Now, confirmed credential compromises trigger immediate containment before an analyst even reviews the incident.

Example 4: Agentic Investigation and Remediation

An authentication failure alert appears for a user account accessing a sensitive database. An AI agent correlates the failed login with the user's normal behavior patterns, recent permission changes, similar activities across other accounts, and current threat intelligence. The agent determines this represents credential compromise, not user error. It immediately disables the account, terminates active sessions, blocks the source IP, and documents its reasoning—closing the loop before an analyst begins their workday.

How We Designed for This Future

Building the autonomous SOC required creating technology that works fundamentally different from traditional security platforms.

A True Platform, Not a Product Suite

Cortex XSIAM unifies your security data in one AI-driven platform, delivering the capabilities of six products in a single solution to detect and stop attacks faster than ever before. Extended detection and response (XDR) eliminates threats targeting endpoints, networks, and clouds before damage occurs. SIEM provides comprehensive visibility and investigation capabilities across your entire security infrastructure. Security orchestration, automation, and response (SOAR) eliminates manual work through intelligent playbooks and closed-loop remediation. Cloud detection and response (CDR) secures your cloud workloads and identifies misconfigurations that create exposure. Exposure management discovers and prioritizes external-facing risks attackers see first. Identity threat detection and response (ITDR) prevents credential theft and stops lateral movement.

One user interface. One backend. One data model. No console switching. No API stitching. No data synchronization delays.

When your endpoint agent detects suspicious process execution, that data immediately correlates with network traffic analysis, cloud API calls, identity authentications, and vulnerability context. The correlation happens automatically because everything operates on unified data.

Your team onboards new data sources in a few clicks instead of requiring custom engineering. Detection logic automatically applies to new data. And, playbooks execute across your entire security stack without brittle integrations breaking.

AI Models Built for Security Outcomes

Generic machine learning (ML) applied to security data produces generic results. We built over 2,600 ML models specifically designed to identify attacks.

These models detect attacker techniques through behavioral analytics and pattern recognition. They identify credential theft and brute force attack attempts. And, they recognize lateral movement and privilege escalation behaviors, as well as detect data exfiltration activities and insider threats.

The platform continuously updates detection models with threat intelligence gathered from over 80,000 enterprise customers globally, including over 30 million new malware samples and 500 billion events daily, translating insights into protections that are automatically distributed across all XSIAM deployments.

When suspicious activity emerges, the models automatically stitch related alerts into complete incidents using technique-based intelligence. They enrich investigations with relevant context. An AI-driven [SmartScoring](#) system prioritizes incidents based on overall risk and business impact, not just technical indicators.

Your analysts receive high-confidence incidents that show complete attack timelines instead of drowning in low-confidence alerts that require manual correlation.

Opinionated Content Delivered Continuously

Traditional vendors provide blank canvas platforms expecting you to build your own security program. You write detection rules, create response playbooks, and develop analytics, maintaining everything as threats evolve.

We believe your security vendor should have an opinion. With over two decades defending enterprises, we know what good security operations look like. You should embed that knowledge into your platform.

Cortex XSIAM delivers opinionated content automatically. It provides detection logic for common attack techniques, response playbooks for standard incident types, analytics for typical security use cases, and dashboards for normal operational visibility.

This content updates continuously based on [Unit 42 Threat Intelligence and incident response](#) engagements. New attack techniques identified in the wild automatically become detections in your environment. You benefit from collective defense intelligence without writing custom rules.

Your team focuses on the security challenges that are unique to your business instead of rebuilding the general risks that every organization faces.

Transformation Your Organization Needs

Moving from today's fragmented security operations to an autonomous SOC is not a technology replacement project. It is a strategic transformation of how your organization approaches defense.

Measure Differently

Redefine how you measure security operations success. Traditional metrics optimize for activity. The modern SOC needs to measure for security quality, including measuring incident recurrence rates that show whether you are eliminating root causes. Track security posture improvement through metrics like reduction in exploitable vulnerabilities and shrinking the attack surface. Measure analyst effectiveness through strategic outcomes like proactive threat discoveries, not reactive alert processing.

These measurements align security operations with business objectives. They demonstrate value in terms executives understand: risk reduction, operational efficiency, and strategic capability development.

Invest in Strategic Skills

SOC transformation also includes understanding the changes in what capabilities your security team needs. When your team spends less time administering tools, they have more time for strategic security work. Your team's role shifts from alert processors to threat hunters, from tool administrators to security architects, from reactive responders to proactive defenders.

Investing in strategic skills solves your talent shortage. Analysts must stay engaged and develop transferable expertise instead of burning out on repetitive work. Your organization builds institutional knowledge that compounds over time instead of always retraining replacements.

To enhance your team's skills, develop the following higher-level capabilities through strategic engagements.

Tabletop Exercises: Build Business Smarts and Strategic Thinking

Your team runs a high-pressure cyber crisis simulation based on real attacks. This forces them to work with executives, weigh the business risk of decisions (like paying a ransom), and test your entire strategic response, not just the technical steps.

Purple Team Exercises: Grow Architectural and Threat Modeling Skills

This is a hands-on workshop where your defensive team works with our ethical hackers. They learn exactly how the bad guys attack your specific systems, which directly improves your team's understanding of security architecture and how to model and defeat threats.

Cyber Risk Assessment: Sharpen Risk Analysis and Threat Intelligence

We assess your defenses against the threats most likely to hit you. Your analysts learn to connect world-class threat intelligence from Unit 42 to your environment, enabling them to precisely calculate business risk and prioritize fixes based on what matters.

Security Program Design: Teach Strategic Vision and Architecture

We help you build the blueprints for your entire security program. Your team gains the experience they need to think like a CISO—defining a strategic plan and security architecture roadmap that ties directly to your company's long-term goals.

After transforming to autonomous security operations, organization operating improvements don't take long to realize.

Your analysts start each day by reviewing a prioritized list of incidents that require human judgment. Routine threats are contained automatically overnight. Complex investigations arrive with complete timelines already reconstructed. Business impact assessments are clear. Recommended actions are specific.

Your team spends most of their time on strategic security work. Proactive threat hunting discovers attacks before they cause damage. Architectural security reviews improve your defensive posture. Collaboration with business units aligns security strategy with organizational objectives.

The Choice You Face

The only question is whether you drive transformation proactively or experience it reactively after your next incident exposes current model limitations.

Organizations that continue operating legacy SOC models face escalating costs, increasing risk, and declining effectiveness. The fragmented, reactive approach is not sustainable.

The autonomous SOC offers a different path. Unified data eliminates correlation gaps. Security-specific intelligence detects threats that manual analysis misses. Outcome-driven automation frees your team to focus on strategic defense. Your organization becomes more secure each day instead of simply processing more alerts.

This transformation requires commitment. Technology alone is not enough. Success demands rethinking security operations from first principles, measuring outcomes instead of activity, and building organizational capabilities for the autonomous era.

Organizations that make this commitment gain decisive advantages, lower operational costs, reduced security risk, and a strategic security capability that enables rather than constrains business objectives. They also gain a security team that grows in expertise instead of burning out on reactive work.

Your next step is not selecting technology but committing to transformation.

Cortex XSIAM provides the unified architecture, security-specific intelligence, and outcome-driven automation that make autonomous security operations possible. The vision and commitment must come from you.

This future is available today for organizations ready to transform. The question is whether your organization will lead this transformation or follow after others demonstrate the advantages.

The choice is yours.

To take the next step in your SOC transformation, read our [definitive guide](#), and see how Cortex XSIAM in action in our product [tour](#).

About Cortex XSIAM

Cortex XSIAM is the AI-driven security operations platform for the modern SOC, harnessing the power of AI to simplify security operations, stop threats at scale, and accelerate incident remediation.

XSIAM unifies best-in-class security operations functions, including EDR, XDR, SOAR, ASM, UEBA, TIP, and SIEM. XSIAM centralizes all of your security data and uses machine learning data models designed specifically for security. With XSIAM, automate data integration, analysis, and response actions, enabling analysts to focus on the incidents that matter.

To learn more about Cortex XSIAM, visit www.paloaltonetworks.com/cortex/cortex-xsiam.

Quantified Benefits: Forrester® TEI™ Study

A Forrester Total Economic Impact™ Study examined organizations that unified their security operations and found measurable outcomes over three years:

Financial Impact

- **257%** 3-year return on investment (ROI)
- **\$5.6M** net present value
- **<6 months** payback period
- **\$7.7M** in total benefits vs. **\$2.2M** in costs

Operational Improvements

- **85%** reduction in alerts requiring Tier 1 attention
- **70%** reduction in incidents requiring investigation
- **85%** reduction in mean time to remediation (MTTR)
- **60%** improvement in security posture (valued at \$2.2M)

Cost Savings

- **\$3.1M** saved from eliminating 20+ legacy tools
- **\$930K** value from improved triage efficiency
- **\$1.5M** value from enhanced case management

Source: [The Total Economic Impact™ Of Palo Alto Networks Cortex XSIAM](#), a commissioned study conducted by Forrester Consulting on behalf of Palo Alto Networks, October 27, 2025. Results are for a composite organization based on interviewed customers.



3000 Tannery Way
Santa Clara, CA 95054
Main: +1.408.753.4000
Sales: +1.866.320.4788
Support: +1.866.898.9087
www.paloaltonetworks.com

© 2025 Palo Alto Networks, Inc. A list of our trademarks in the United States and other jurisdictions can be found at <https://www.paloaltonetworks.com/company/trademarks.html>. All other marks mentioned herein may be trademarks of their respective companies.
cortex_wp_your-security-future_112525